

Revision Plan for SIGCOMM'24 Paper #293

Dear Prof. Praveen Kumar,

Thank you for shepherding our paper (#293 *PPT: A Pragmatic Transport for Datacenters*). We carefully went through the reviews and outlined our revision plan.

Please let us know if you have any feedback. Thank you for your time and help to improve the quality of our paper!

Best Regards,
Authors of SIGCOMM'24 Paper #293

I Top-level Concerns:

Concern 1: *Disconnect between the proposed goals vs the evaluation. The paper pitches itself to deliver comparable performance to proactive transports while being readily deployable. But the evaluations do not necessarily back this up.*

Response 1: Thanks for the reviewer's comments. We aim to explore a new transport that can achieve comparable performance to proactive transports while being readily deployable. To this point, we start with DCTCP and use LCP to gracefully utilize the available bandwidth. Moreover, we complement its design with a buffer-aware flow scheduling to optimize small flow's performance. We then show that PPT can achieve performance comparable to proactive transports by overall average, small flow average/ tail, and large flow average FCT. We proved the readily deployable nature of PPT by implementing PPT based on Linux-kernel and low CPU overhead. The above performance metrics can back up our design goals.

Concern 2: *There's lacking discussion on other related work such as PCC Proteus.*

Response 2: Thanks for the reviewer's comments. We will add a discussion of PCC Proteus and other related work in Appendix B. However, note that PCC Proteus and PPT are two lines of work, so we have yet to discuss it. PCC Proteus divides flows into primary and scavenger flows according to different application requirements and proactively reduces scavenger flow bandwidth to minimize the impact on the primary flow, thus improving the overall user experience. By contrast, PPT is insensitive to the application requirements and treats all flows equally. PPT uses a dual-loop rate control design to utilize the available bandwidth gracefully and further complements its design with a buffer-aware flow scheduling to optimize small flow's performance.

Concern 3: *The paper needs to discuss more insights on how the system is able to deliver the performance benefits.*

Response 3: Thanks for the reviewer's comments. We will provide a more detailed discussion of how the system can deliver the performance benefits. PPT uses LCP to gracefully utilize the spare bandwidth left

by DCTCP in the slow start and queue buildup phases. In particular, we will add further discussion for the queue buildup phase to show more details in Section 2.3. Furthermore, we complement buffer-aware flow scheduling to optimize the performance of small flows. We will also add a discussion of buffer-aware flow scheduling in Appendix Section 3.1.

Concern 4: *Other claims such as "integration with other transports is easy" need to be further justified.*

Response 4: Thanks for the reviewer's comments. Actually, we have not claimed that "integration with other transports is easy." We claim that the design of PPT can be integrated with delay-based transport. To this point, we implemented a delay-based transport conceptually equivalent to Swift in the ns3 simulator. We compared this variant with the original delay-based transport and plotted the result in Figure 26. We find that when incorporating PPT's design with the original delay-based transport, the overall average FCT, the average/tail FCT of small flows, and the average FCT of large flows can be reduced by 16.7%, 56.5%/72.1%, and 11%, respectively. To show more details, we plan to add a more detailed description of this variant and discuss how delay-based transport can benefit from the design of the PPT in Appendix C.3.

Concern 5: *Missing details on evaluation setup and parameters.*

Response 5: Thanks for the reviewer's comments. We will complete the missing details about the experimental setup and parameters in the form of tables or footnotes.

II Other Concerns

1 Reviewer #293A

A.1: Reviewer Comment: *The number of used queues could be clarified: p. 6 mentions two such queues (a high-priority and a low-priority queue with a different λ . But then it mentions that when queueing the opportunistic (presumably low-priority) packets, this may hurt normal (presumably high-priority) packets, suggesting they share the queue. Then, p. 8 mentions 8 priorities, such that "switches can use strict priority to dequeue packets", implicitly pointing to 8 different queues. Then, p.9 mentions 2 queues again.*

Response A.1: Thanks for the reviewer's comments. PPT uses eight strict priorities in the switch. To prevent HCP traffic harmed by LCP, priorities 0~3 are used to transmit HCP packets, and priorities 4~7 are used to transmit LCP packets. Therefore, the *high-priority* and *low-priority* mentioned in p.6 and p.9 refer to priorities 0~3 and 4~7 in the switch, respectively. To avoid ambiguity, we will include footnotes in Section 3.2 for explanation.

A.2: Reviewer Comment: *In addition, PPT seems to couple each LCP flow with an HCP flow, by defining the window size of the LCP flow using the window size of the associated HCP flow. But what if there are no HCP flows now? Are LCP flows stuck? Or more generally, what if the numbers of flows are not equal? Or what if they are equal but the flows are destined to different destinations? How does this coupling help?*

Response A.2: Thanks for the reviewer's comments. Although very sorry, the reviewer seems to have misinterpreted PPT's strategy. At the beginning of Section 3, we mentioned that "LCP sends opportunistic packets from the tail end." Therefore, both HCP and LCP send data in the same flow. However, many of the review's comments were based on LCP and HCP transmitting data from different flows, which made it difficult for us to answer.

2 Reviewer #293B

B.1: Reviewer Comment: *During LCP loop initialization for case 1, it is unclear how PPT computes BDP in the first RTT. As you mention, DCTCP is still probing for available bandwidth.*

Response B.1: Thanks for the reviewer’s comments. Since the NIC rate can be read via system call and the datacenter network topology is relatively fixed, we assume the RTT is known a priori. By NIC rate times RTT, we assume the host knows the BDP in advance.

B.2: Reviewer Comment: *I wonder if tracking α_{min} can lead to underutilization due to network dynamics. Suppose we have a long flow which experiences transient congestion at the beginning (say due to an incast); the transient congestion causes α_{min} to be > 0.5 for the flow. After the transient congestion subsides, the large flow is unable to utilize any spare bandwidth arising from Case 2.*

Response B.2: Thanks for the reviewer’s comments. During the process of congestion subsiding, the switch queue length is reduced, the percentage of ECN marked packets for this long flow decreases, and the value of α . As the congestion subsides completely, the proportion of ECN-marked packets should be 0. PPT initializes LCP for the long flow whenever α takes the minimum value. At this point, α must be less than 0.5, which makes this large flow utilize the spare bandwidth.

B.3: Reviewer Comment: *It is unclear how PPT deals with fairness as flows arrive and leave. Suppose two large flows sharing a common bottleneck link arrive one after the other. The latter flow might see much lower W_{max} (50%?) than the earlier flow. In every LCP loop initialization (case 2), the latter flow would compute a lower value of initial congestion window and send less packets via LCP compared to the earlier flow.*

Response B.3: Thanks for the reviewer’s comments.

B.4: Reviewer Comment: *As the paper targets the issue of underutilizing high datacenter bandwidth, I wonder if some of the simulations could have been done at higher line rates (400+ Gbps) to demonstrate the problem more clearly.*

Response B.4: Thanks for the reviewer’s comments. We will add a new simulation experiment with topology consisting of 144 servers, 9 leaf switches, and 4 spine switches, with the host and core links operated at 400 and 1000Gbps, respectively. We will plot the results in the appendix.

B.5: Reviewer Comment: *What are the overheads of PPT?*

Response B.5: Thanks for the reviewer’s comments. Actually, we measure the kernel space CPU overhead of PPT and DCTCP in our testbed. We plot the results in Figure 21. For more details, please read Appendix C.1.

3 Reviewer #293C

C.1: Reviewer Comment: *It would be particularly useful to have an intuitive and straightforward example scenario where LCP can discover spare bandwidth other than during the slow start phase.*

Response C.1: Thanks for the reviewer’s comments. Indeed, in section 2.3, we mention that DCTCP marks ECN at the switch for arriving packets if queue occupancy exceeds a threshold K , and the sender cuts the window based on the fraction of ECN marked ACKs. So, when there are multiple concurrent flows, DCTCP may mark ECN for packets from many flows, which may cut windows simultaneously, thus causing a sudden drain on the switch buffer and leaving bandwidth underutilized. To show this point, We run ns-3 simulations

with two senders and one receiver sharing the bottleneck. We sample the bottleneck link utilization every 100us for 10ms when DCTCP enters a steady state. We plot the results in Figure 1, showing that nearly half of the bandwidth is underutilized.

C.2: Reviewer Comment: *Figure 13 doesn't add up by itself. RC3 has the highest average FCT for small and large flows. However, its overall average is not the highest. Given that large and small flows account for 13% and 87% respectively, the overall average should be around $0.13 * 39.63 + 0.87 * 0.77 = 5.8$? Or am I missing something?*

Response C.2: Thanks for the reviewer's comments. After carefully checking our data processing code, we found that under Datamining workloads, RC3 would result in a few flows not completing due to many packet losses. Our code incorrectly used the program runtime as the flow completion time when calculating the overall average FCT, resulting in the actual computed result being larger. We will fix the bug and update the correct result.

C.3: Reviewer Comment: *Page 12, can you give an intuition why limiting RC3's low priority queue doesn't help? It kind of contradicts the argument that aggressive low priority packets hinder high-priority packets being a problem if the high-priority queue has enough buffers?*

Response C.3: Thanks for the reviewer's comments. Although limiting the RC3 low-priority queue can prevent the LCP from sending too many opportunity packets, the following reasons still lead to low performance. First, during the slow-start phase of the flow, limiting the buffer LCP could use will leave a large amount of bandwidth underutilized. Second, RC3 keeps the LCP open instead of intermittent initialization as PPT, which causes more severe congestion during traffic bursts. Last but not least, RC3's lack of using in-network priority causes small flow packets to be queued after LCP packets, increasing small flow latency. We will add the above analyses to our paper.

4 Reviewer #293D

D.1: Reviewer Comment: *I'm not sure if I fully understand some of the experimental results. To give an example, why is Homa's performance so poor for the testbed incast experiment? You mention that PPT is effective because it can recover quickly but why can't Homa?*

Response D.1: Thanks for the reviewer's comments. As mentioned in Section 6.1.2, Homa-Linux sends BDP-sized unscheduled packets for each arriving flow, and excessive opportunity packets cause heavier congestion in incast scenarios. Furthermore, Homa-Linux is connectionless, so packet loss recovery is achieved through retransmission after timeout. Large packet loss in Incast scenarios further impairs its performance. We will add more detailed analyses in the paper.

5 Reviewer #293E

E.1: Reviewer Comment: *In production systems, priority queues are scarce resource. Using only two HW queues for one application will be a luxury. How does PTP perform with 2 prio queues (one high, one low) compared to existing production solutions like Swift, HPCC, DCTCP?*

Response E.1: Thanks for the reviewer's comments. In our paper, we construct a stricter variant of PPT that assigns all packets the same priority. We plot the results in Figure 19. Even under these conditions, we find that PPT still achieves the optimal overall average FCT among NDP, Homa, Aeolus, DCTCP, and RC3.

E.2: Reviewer Comment: *Even the 'large-scale' simulations cover only 144 servers w/ 40 100G as core link speeds. This resembles DC topology of more a decade ago. With the trend of AI clusters growing to have 10s of Ks of GPUs, I believe at least 2K or more endpoints must be simulated to see the real impact on shallow switch.*

Response E.2 Thanks for the reviewer's comments. We will add a large-scale simulation experiment with 2k nodes to the appendix.

References

- [1] V. Dukic, S. A. Jyothi, B. Karlas, M. Owaida, C. Zhang, and A. Singla, "Is advance knowledge of flow sizes a plausible assumption?" in *Proc. of USENIX NSDI*, 2019.
- [2] A. Rashelbach, O. Rottenstreich, and M. Silberstein, "A computational approach to packet classification," in *Proc. of ACM SIGCOMM*, 2020.